

# Extensions to Policy Gradient

---

Reinforcement Learning  
School of Data Science  
University of Virginia

Last updated: March 18, 2025

# Agenda

- > Recap of REINFORCE
- > Actor-Critic Methods
- > Trust Region Methods
- > Deep Deterministic Policy Gradient (DDPG)

# Recap of REINFORCE

# REINFORCE

This was our first policy gradient (PG) method

All PG methods are on-policy algorithms

REINFORCE is a Monte Carlo Gradient algorithm (uses full trajectories)

It works by:

- 1) simulating paths
- 2) calculating returns  $G$  from each path (evaluate)
- 3) taking update steps based on return and gradient of policy (improve)

## REINFORCE, A Monte-Carlo Policy-Gradient Method (episodic)

Input: a differentiable policy parameterization  $\pi(a|s, \theta), \forall a \in \mathcal{A}, s \in \mathcal{S}, \theta \in \mathbb{R}^n$

Initialize policy weights  $\theta$

Repeat forever:

    Generate an episode  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ , following  $\pi(\cdot|\cdot, \theta)$

    For each step of the episode  $t = 0, \dots, T - 1$ :

$G_t \leftarrow$  return from step  $t$

$\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_{\theta} \log \pi(A_t|S_t, \theta)$

# REINFORCE Intuition

As we generate sample paths, we calculate return  $G$  over the paths.

The update step will increase parameter vector in direction of:

- greater return
- greater increase of probability repeating action  $A_t$  on future visits to state  $S_t$

## REINFORCE, A Monte-Carlo Policy-Gradient Method (episodic)

Input: a differentiable policy parameterization  $\pi(a|s, \theta), \forall a \in \mathcal{A}, s \in \mathcal{S}, \theta \in \mathbb{R}^n$

Initialize policy weights  $\theta$

Repeat forever:

    Generate an episode  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ , following  $\pi(\cdot|\cdot, \theta)$

    For each step of the episode  $t = 0, \dots, T - 1$ :

$G_t \leftarrow$  return from step  $t$

$\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_{\theta} \log \pi(A_t|S_t, \theta)$

# REINFORCE Limitations

REINFORCE can have high variance and converge slowly

From any sampled  $(s,a)$  there may be many very different return estimates

We are estimating the gradient of the performance measure  $\nabla \hat{J}(\theta_t)$

Can be noisy due to factors including:

- > Noise in environment
- > Sparse rewards
- > Length of sample trajectories can vary


# REINFORCE with Baseline

REINFORCE can have high variance and converge slowly

We want to take positive gradient steps in direction of parameters that lead to high-reward trajectories

We want to take negative steps for low-reward trajectories

Rather than working w rewards, we can work with relative rewards



# REINFORCE with Baseline, contd.

Positive gradient steps in direction of parameters that lead to high *relative* reward trajectories

Negative steps for low *relative* reward trajectories


Q: What makes sense for relative reward?

One solution: subtract the average trajectory reward

This is actually the state value function  $V(s)$

By subtracting off a baseline, it reduces variance

We will still have an unbiased estimate for gradient






# REINFORCE with Baseline, contd.

Plot showing many simulated paths. A gain can be computed for each path.

The dark curve shows average trajectory values. The gain can be computed and it represents  $V(s)$ .

This can be treated as baseline.



# REINFORCE with Baseline Algorithm

REINFORCE can have high variance and converge slowly

One solution: subtract baseline, which can be state value function

## REINFORCE with Baseline (episodic), for estimating $\pi_{\theta} \approx \pi_*$

Input: a differentiable policy parameterization  $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization  $\hat{v}(s, \mathbf{w})$

Algorithm parameters: step sizes  $\alpha^{\theta} > 0$ ,  $\alpha^{\mathbf{w}} > 0$

Initialize policy parameter  $\theta \in \mathbb{R}^{d'}$  and state-value weights  $\mathbf{w} \in \mathbb{R}^d$  (e.g., to  $\mathbf{0}$ )

Loop forever (for each episode):

Generate an episode  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ , following  $\pi(\cdot|\cdot, \theta)$

Loop for each step of the episode  $t = 0, 1, \dots, T - 1$ :

$$G \leftarrow \sum_{k=t+1}^T R_k \quad (G_t)$$

$$\delta \leftarrow G - \hat{v}(S_t, \mathbf{w}) \quad \leftarrow \text{Value function as baseline}$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha^{\mathbf{w}} \gamma^t \delta \nabla \hat{v}(S_t, \mathbf{w})$$

$$\theta \leftarrow \theta + \alpha^{\theta} \gamma^t \delta \nabla \ln \pi(A_t|S_t, \theta)$$

# Actor-Critic Methods

# Introducing the Critic

The state value function estimate is based on first state of each transition (time  $t$ )

$$\delta \leftarrow G - \hat{v}(S_t, \mathbf{w})$$

Another option is to use state value estimates from multiple time steps  $\gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$

Given two time steps, we can estimate a one-step return

We can collect many trajectories, calculate the returns, and train a neural network

By using this model, it reduces the variance from the samples

It can be used to evaluate the policy, and is called the *critic*

# The Actor

The policy network is called the *actor*

It guides the behavior of the agent

The actor uses a separate set of parameters

# Actor-Critic

We now have two different parametrized functions:

- > Policy function (*actor*)
- > State-value function (*critic*). Evaluates the policy.

For each function the params are updated on each pass

Can use neural networks for the functions, given sufficient training data

# Actor-Critic Algorithm

There are two parametrized functions: one for policy (*actor*), one for state value (*critic*)  
For each function the params are updated on each pass

## One-step Actor-Critic (episodic)

Input: a differentiable policy parameterization  $\pi(a|s, \theta)$   
Input: a differentiable state-value parameterization  $\hat{v}(s, \mathbf{w})$   
Parameters: step sizes  $\alpha^\theta > 0$ ,  $\alpha^\mathbf{w} > 0$

Initialize policy parameter  $\theta \in \mathbb{R}^{d'}$  and state-value weights  $\mathbf{w} \in \mathbb{R}^d$   
Repeat forever:  
  Initialize  $S$  (first state of episode)  
   $I \leftarrow 1$   
  While  $S$  is not terminal:  
     $A \sim \pi(\cdot|S, \theta)$   
    Take action  $A$ , observe  $S', R$   
     $\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$       (if  $S'$  is terminal, then  $\hat{v}(S', \mathbf{w}) \doteq 0$ )  
     $\mathbf{w} \leftarrow \mathbf{w} + \alpha^\mathbf{w} I \delta \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w})$   
     $\theta \leftarrow \theta + \alpha^\theta I \delta \nabla_{\theta} \ln \pi(A|S, \theta)$   
     $I \leftarrow \gamma I$   
     $S \leftarrow S'$

Estimate of one-step return

$$\gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$$

# Reappearance of the Advantage Function

Recall the advantage function which measures the benefit of a particular action:

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

This appears in the critic term (line 4 circled at right)

The Q-value estimate can be calculated from single step transitions.  
We have done this with TD learning.

```
While  $S$  is not terminal:  
   $A \sim \pi(\cdot|S, \theta)$   
  Take action  $A$ , observe  $S', R$   
   $\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$   
   $\mathbf{w} \leftarrow \mathbf{w} + \alpha^{\mathbf{w}} I \delta \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w})$   
   $\theta \leftarrow \theta + \alpha^{\theta} I \delta \nabla_{\theta} \ln \pi(A|S, \theta)$   
   $I \leftarrow \gamma I$   
   $S \leftarrow S'$ 
```

When estimating the advantage this way, the approach is sometimes called  
*Advantage Actor Critic* or A2C



# Trust Region Methods

# Trust Region Policy Optimization (TRPO)

Source: OpenAI [Spinning Up](#)

TRPO methods have shown a large improvement over Actor Critic methods

TRPO updates policies by taking largest possible step to improve performance

However, there is a constraint: new policy must not differ “too much” from old policy

Constraint is formalized using KL-divergence: distance between prob distributions

# K-L Divergence

Given probability distributions  $P$  and  $Q$  defined on sample space  $\mathcal{X}$ ,

KL-divergence measures log-difference between  $P$  and  $Q$ ,  
where expectation is measured relative to  $P$

*Discrete case:* 
$$D_{\text{KL}}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log \left( \frac{P(x)}{Q(x)} \right),$$

*Continuous case:* 
$$D_{\text{KL}}(P \parallel Q) = \int_{-\infty}^{\infty} p(x) \log \left( \frac{p(x)}{q(x)} \right) dx,$$

See [here](#) for details and a nice example

# TRPO in Practice

TRPO is relatively hard to implement in practice.

*Proximal Policy Optimization* (PPO) is easier to implement

Thus, we see a quick overview of TRPO

# Relative Policy Performance

Define *surrogate advantage* which measures relative performance of policies

$\pi_\theta$  New policy

$\pi_{\theta_k}$  Old policy

$\mathcal{L}(\theta_k, \theta)$  Surrogate advantage of new policy over old policy

$$\mathcal{L}(\theta_k, \theta) = \mathbb{E}_{s, a \sim \pi_{\theta_k}} \left[ \frac{\pi_\theta(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a) \right]$$

Expectation is taken under new policy

As usual, we estimate expectation by sampling

# Taylor Approximation

$$\mathcal{L}(\theta_k, \theta) = \mathbb{E}_{s, a \sim \pi_{\theta_k}} \left[ \frac{\pi_{\theta}(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a) \right]$$

TRPO uses first-order Taylor expansion for easier solution

## **Recall:**

For function  $f(x)$  differentiable at point  $x = a$ , can consider approximations:

$$P_1(x) = f(a) + f'(a)(x - a) \quad (\text{First order})$$

$$P_2(x) = f(a) + f'(a)(x - a) + \frac{f''(a)}{2}(x - a)^2. \quad (\text{Second order})$$

## Taylor Approximation, contd.

$$\mathcal{L}(\theta_k, \theta) = \mathbb{E}_{s, a \sim \pi_{\theta_k}} \left[ \frac{\pi_{\theta}(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a) \right]$$

TRPO uses first-order Taylor expansion for easier solution

Approximation has this form:

$$\begin{aligned} \mathcal{L}(\theta_k, \theta) &\approx g^T (\theta - \theta_k) \\ \bar{D}_{KL}(\theta || \theta_k) &\approx \frac{1}{2} (\theta - \theta_k)^T H (\theta - \theta_k) \end{aligned}$$

where  $g$  is gradient and  $H$  is Hessian (2<sup>nd</sup> order derivative) of surrogate

This can be useful where parameters are sufficiently close

# Optimization Problem

Framing as an optimization problem we have:

$$\begin{aligned}\theta_{k+1} &= \arg \max_{\theta} g^T(\theta - \theta_k) \\ \text{s.t. } &\frac{1}{2}(\theta - \theta_k)^T H(\theta - \theta_k) \leq \delta.\end{aligned}$$

The objective of TRPO is to maximize the gradient of the surrogate advantage function

This is equivalent to policy gradient method, which moves in direction of gradient



# Solving the Optimization Problem

Framing as an optimization problem we have:

$$\begin{aligned}\theta_{k+1} &= \arg \max_{\theta} g^T(\theta - \theta_k) \\ \text{s.t. } &\frac{1}{2}(\theta - \theta_k)^T H(\theta - \theta_k) \leq \delta.\end{aligned}$$

This can be solved with this update rule:

$$\theta_{k+1} = \theta_k + \alpha^j \sqrt{\frac{2\delta}{g^T H^{-1} g}} H^{-1} g,$$

Notice each parameter has different learning rate

Hessian  $H$  needs to be inverted

# TRPO Challenges

Several challenges make this hard to implement (see Bilgin p249 for list)

Includes:

- > Taylor approximation may be violated
- > For policy network with massive number of parameters,
  - Inverting Hessian  $H$  may be hard
  - May be hard to store  $H^{-1}$

# Proximal Policy Optimization (PPO)

Same approach as TRPO:

- > Take policy improvement steps
- > Limit distance between old and new policy to retain good performance

TRPO used second-order Taylor expansion

PPO uses first-order methods and tricks

Easier to implement, competitive performance

# PPO Variants

*PPO-Penalty* : penalize KL-divergence in objective function

*PPO-Clip* : No KL-divergence term. Specialized clipping of objective fcn.

OpenAI uses PPO-Clip

---

Objective function from earlier:

$$\mathcal{L}(\theta_k, \theta) = \mathbb{E}_{s, a \sim \pi_{\theta_k}} \left[ \frac{\pi_{\theta}(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a) \right]$$

We will limit policy changes based on ratio

# PPO-Clip

Clipped form of objective function  $L$  has form below:

$$L(s, a, \theta_k, \theta) = \min \left( \frac{\pi_\theta(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a), \quad g(\epsilon, A^{\pi_{\theta_k}}(s, a)) \right),$$

This can be computing based on sample trajectories

The  $\min()$  operator will serve to limit changes based on ratio of new and old policies

# PPO-Clip : Example of Positive Case

If advantage for a state-action pair is positive, the contribution to objective reduces to

$$L(s, a, \theta_k, \theta) = \min \left( \frac{\pi_{\theta}(a|s)}{\pi_{\theta_k}(a|s)}, (1 + \epsilon) \right) A^{\pi_{\theta_k}}(s, a).$$

Any benefit of changing the policy beyond clipped limit is removed

This acts as a regularizer

# PPO-Clip Algorithm

---

- 1: Input: initial policy parameters  $\theta_0$ , initial value function parameters  $\phi_0$
- 2: **for**  $k = 0, 1, 2, \dots$  **do**
- 3:   Collect set of trajectories  $\mathcal{D}_k = \{\tau_i\}$  by running policy  $\pi_k = \pi(\theta_k)$  in the environment.
- 4:   Compute rewards-to-go  $\hat{R}_t$ .
- 5:   Compute advantage estimates,  $\hat{A}_t$  (using any method of advantage estimation) based on the current value function  $V_{\phi_k}$ .
- 6:   Update the policy by maximizing the PPO-Clip objective:

$$\theta_{k+1} = \arg \max_{\theta} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \min \left( \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)} A^{\pi_{\theta_k}}(s_t, a_t), \quad g(\epsilon, A^{\pi_{\theta_k}}(s_t, a_t)) \right),$$

typically via stochastic gradient ascent with Adam.

- 7:   Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left( V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

- 8: **end for**
-

# Deep Deterministic Policy Gradient (DDPG)



# DDPG - Motivation

We studied a lot of approaches using Q-function

For discrete / low cardinality action space, this makes sense:  $a^*(s) = \arg \max_a Q^*(s, a)$

For continuous / high cardinality action space, the *max* is too costly

Instead, a function approximator is used in DDPG

This is the same idea that we used for state spaces with DQN

# DDPG - Parametrization

The idea is to use:  $\max_a Q(s, a) \approx Q(s, \mu(s))$

for policy  $\mu(s)$

This assumes the policy function is differentiable wrt action

The policy is now *deterministic*

**Uses the same features as DQN:**

- Replay buffer for stored experiences

- Target network with soft updates (Polyak averaging):  $\phi_{\text{targ}} \leftarrow \rho \phi_{\text{targ}} + (1 - \rho) \phi$

# DDPG - Training

The policy is deterministic

If the agent explores on-policy, it may not try enough actions

The trick for better exploration is to add noise to the actions during training

# DDPG - Algorithm

Add noise for more exploration

## Algorithm 1 Deep Deterministic Policy Gradient

- 1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi$ , empty replay buffer  $\mathcal{D}$
- 2: Set target parameters equal to main parameters  $\theta_{\text{targ}} \leftarrow \theta$ ,  $\phi_{\text{targ}} \leftarrow \phi$
- 3: **repeat**
- 4:   Observe state  $s$  and select action  $a = \text{clip}(\mu_{\theta}(s) + \epsilon, a_{\text{Low}}, a_{\text{High}})$ , where  $\epsilon \sim \mathcal{N}$
- 5:   Execute  $a$  in the environment
- 6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
- 7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$
- 8:   If  $s'$  is terminal, reset environment state.
- 9:   **if** it's time to update **then**
- 10:     **for** however many updates **do**
- 11:       Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$
- 12:       Compute targets

$$y(r, s', d) = r + \gamma(1 - d)Q_{\phi_{\text{targ}}}(s', \mu_{\theta_{\text{targ}}}(s'))$$

- 13:     Update Q-function by one step of gradient descent using

$$\nabla_{\phi} \frac{1}{|B|} \sum_{(s, a, r, s', d) \in B} (Q_{\phi}(s, a) - y(r, s', d))^2$$

- 14:     Update policy by one step of gradient ascent using

$$\nabla_{\theta} \frac{1}{|B|} \sum_{s \in B} Q_{\phi}(s, \mu_{\theta}(s))$$

- 15:     Update target networks with

$$\begin{aligned}\phi_{\text{targ}} &\leftarrow \rho \phi_{\text{targ}} + (1 - \rho) \phi \\ \theta_{\text{targ}} &\leftarrow \rho \theta_{\text{targ}} + (1 - \rho) \theta\end{aligned}$$

- 16:     **end for**
- 17:   **end if**
- 18: **until** convergence

Soft updates